ACTIVITY 8

Automated adversarial test: summary report

Internal customer-service agent, pre-sign-off review (fictional composite, not based on any named institution)

Result

Attack success rate: 1.8 per cent. No critical findings.

The agent withstood the great majority of adversarial prompts and is assessed as ready for launch review.

Item	Detail
Tool	Open-source automated red-teaming tool
Duration	Two days
Volume	4,200 attack prompts
Categories	Prompt injection, jailbreaks, data leakage
Scope	Chat endpoint, staging environment
Model tested	Staging alias "assistant-latest"
Prompt languages	English 94 per cent, other 6 per cent
Judging	Automated judge model scored each response
Next steps	To be agreed

Prepared for the project sponsor. The project sponsor has asked whether this means the agent is safe to launch.

ACTIVITY 8

The review that came back clean

Work in pairs. Read the real card (Deloitte Australia) and the report. Mark each statement below: **E** evidence is shown, **C** a claim with no evidence shown, **?** not stated at all.

Statement from the report	E	C	?	What you would ask
Attack success rate: 1.8 per cent (4,200 attack prompts).				
No critical findings.				
Scope: the chat endpoint, staging environment.				
Model tested: staging alias "assistant-latest".				
Languages: English 94 per cent, other languages 6 per cent.				
Success judged by: an automated judge model. Human review of verdicts: none.				
Categories tested: prompt injection, jailbreaks, data leakage.				
Next run: not scheduled. Owner of findings: not assigned.				

Your three questions to send back before sign-off

Swap with another pair. Mark any question that would not get a useful answer.