Handout: Working Definitions for the Week

Terminology · 09:40 – 10:40

The reference for the week. Terms are alphabetical. The right-hand column shows where each term gets its main treatment. Write your own team's variations in the margin: the differences are useful.

Term	Working definition	Where it returns
Agent	A model in a loop that chooses its own steps and tools to pursue a goal	Day 1 autonomy
Agentic RAG	RAG where the agent decides when and what to retrieve, possibly more than once	Day 1 retrieval
AI-BoM	An inventory of everything that determines an agent's behaviour (model, prompts, tools, data sources, policies), versioned so it can be compared over time	Day 3
API	The interface an application uses to call a service, including a model	Throughout
Blast radius	What is affected if a system is wrong, and for how long	Day 1 onwards
Chunk	The piece of a document that is embedded and retrieved	Day 1 retrieval
Compaction	Summarising earlier conversation to free space in the context window	Day 1 memory
Context window	The text the model can consider in one call; not memory	Day 1 memory
Copilot / assistant	A model that helps a person, who decides what happens next	Day 1 autonomy
Data leakage	Sensitive data appearing where it should not, through prompts, memory, retrieval or logs	Day 4
Deterministic	The same input always gives the same output	Day 2
Drift	Behaviour changing without anyone deciding it should, because a model, prompt, tool or data source changed	Day 3
Embedding	A list of numbers representing the meaning of text, so similar meanings are close together	Day 1 retrieval
Excessive agency	Giving an agent more tools, permissions or autonomy than the task needs	Day 2
Gateway	A control point between an agent and what it calls, where requests are checked and logged	Day 2
Grounding	Tying an answer to retrieved sources so it can be checked	Day 1 retrieval
Guardrail	A check on an input, an output or an action, which can be rules, a classifier or a person; it has an error rate	Day 4
Hallucination	A fluent but wrong or invented answer; a property of how models work	Day 1
Harness / orchestrator	The code around the model that runs the loop, holds memory and enforces limits	Day 2
Idempotency	Doing the same operation twice has the same effect as doing it once	Day 3
Inference	Running a model to get an output; where it happens is a residency question	Day 4
Jailbreak	A user deliberately talking a model past its rules	Day 3
MCP	Model Context Protocol: a standard way for agents to discover and call tools and data sources	Day 2
Non-deterministic	The same input can give different outputs	Day 1, Day 2
Prompt	The instructions and input sent to a model, usually a system prompt plus the user's message	Day 1
Prompt injection	Instructions hidden in content the model reads, which the model then follows	Day 3
RAG	Retrieval-augmented generation: look up relevant text first, then give it to the model with the question	Day 1 retrieval
Red teaming	Deliberately attacking your own system to find weaknesses before someone else does	Day 4
Security trimming	Making sure retrieval returns only what the requesting identity could open directly	Day 1 retrieval
Skill	A packaged set of instructions and tools an agent can load; treat it like third-party code	Day 2
Supply chain	The models, prompts, tools and servers a system depends on but did not write	Day 2
Token	The unit a model reads and writes, roughly three-quarters of a word; limits and cost are counted in tokens	Day 1 memory
Tool calling	The model asks for an action; the application runs it and returns the result	Day 1, Day 2
Trace	The step-by-step record of what an agent saw, decided and did	Day 4

Vector store	A database that finds the embeddings nearest to a query	Day 1 retrieval
Workflow	A fixed sequence of steps, set by the builder, where a model fills in some of them	Day 1 autonomy